

Big Data Systems (CS4545/CS6545)  
Winter 2021

# Query Processing

Suprio Ray


University of New Brunswick

# Acknowledgement

---

Thanks to N. Koudas, Manos Papagelis, Jordan Halterman, Ramakrishna and Gehrke, Silberchatz, Korth and Sudarshan for materials in these slides

# Outline

- Quick SQL *review* 
- Query processing overview
- Relational Algebra *review*
- Introduction to Apache Calcite
- Join processing

# Query processing language: SQL (Structured Query Language)

- Declarative
  - Say “what to do” rather than “how to do it”
    - Avoid data-manipulation details needed by procedural languages
  - Database engine figures out “best” way to execute query
    - Called “query optimization”
    - Crucial for performance: “best” can be a million times faster than “worst”
- Data independent
  - Decoupled from underlying data organization
    - Correctness always assured... performance not so much
  - SQL is standard and (nearly) identical among vendors

# SQL Environment

- Catalog

- ➔ – A set of schemas that constitute the description of a database

List of databases					
Name	Owner	Encoding	Collate	Ctype	Access privileges
XDBSYS	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	
__mytest__N1	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	
__xtest__N1	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	
dbgeo	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	
postgres	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	
template0	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	=c/dbuser + dbuser=CTc/dbuser
template1	dbuser	UTF8	en_CA.UTF-8	en_CA.UTF-8	=c/dbuser + dbuser=CTc/dbuser
(7 rows)					

# SQL Environment

- Catalog
- Schema



- The structure that contains descriptions of objects created by a user
- A schema is a collection of related objects, including but not limited to base tables, views, constraints, domains, character sets, triggers and roles

List of relations			
Schema	Name	Type	Owner
public	arealm_merge_ca	table	dbuser
public	arealm_merge_ca_gid_seq	sequence	dbuser
public	arealm_merge_ca_shall	table	dbuser
public	arealm_merge_ca_shqtall	table	dbuser
public	areawater_merge_ca	table	dbuser
public	areawater_merge_ca_gid_seq	sequence	dbuser
public	areawater_merge_ca_shall	table	dbuser
public	areawater_merge_ca_shqtall	table	dbuser
public	edges_merge_ca	table	dbuser
public	edges_merge_ca_gid_seq	sequence	dbuser
public	edges_merge_ca_shall	table	dbuser
public	edges_merge_ca_shqtall	table	dbuser
public	geography_columns	view	dbuser
public	geometry_columns	view	dbuser
public	raster_columns	view	dbuser
public	raster_overviews	view	dbuser
public	spatial_ref_sys	table	dbuser
(17 rows)			

# SQL Environment

- Catalog
- Schema
- Data Definition Language (DDL)



– Commands that **define a database**, including creating, altering, and dropping tables and establishing constraints


```
dbgeo=# \d arealm_merge_ca
Table "public.arealm_merge_ca"
  Column      |      Type      | Modifiers
-----+-----+-----
 gid          | integer        | not null default nextval('arealm_me
 statefp      | character varying(2) |
 countyfp     | character varying(3) |
 ansicode     | character varying(8) |
 areaid       | character varying(22) |
 fullname     | character varying(100) |
 mtfcc        | character varying(5)  |
 aland        | double precision      |
 awater       | double precision      |
 intptlat     | character varying(11) |
 intptlon     | character varying(12) |
 geom         | geometry(Polygon,4326) |
Indexes:
  "arealm_merge_ca_pkey" PRIMARY KEY, btree (gid)
  "arealm_merge_ca_geom_gist" gist (geom)
```

# SQL Environment

- Catalog
- Schema
- Data Definition Language (DDL)
  - Commands that **define a database**, including creating, altering, and dropping tables and establishing constraints
- Data Manipulation Language (DML)
  - Commands that maintain and query a database
  - Commands for **inserting**, **modifying** and **querying** the data in the database



# SQL Environment

- Catalog
- Schema
- Data Definition Language (DDL)
  - Commands that **define a database**, including creating, altering, and dropping tables and establishing constraints
- Data Manipulation Language (DML)
  - Commands that maintain and query a database
  - Commands for **updating**, **inserting**, **modifying** and **querying** the data in the database
- Data Control Language (DCL)
  -  – Commands that **control** a database, including **administering privileges** and committing data

# Basic Single-Table Queries

- `SELECT` `[DISTINCT]` *<column expression list>*  
    `FROM` *<single table>*  
    `[WHERE` *<predicate>*`]`  
    `[GROUP BY` *<column list>*  
        `[HAVING` *<predicate>*`]` `]`  
    `[ORDER BY` *<column list>*`]`
- Simplest version is straightforward
  - Produce all tuples in the table that satisfy the predicate
  - Output the expressions in the `SELECT` list
    - Expression can be a **column reference**, or an **arithmetic expression** over column refs

# Basic Single-Table Queries

- **SELECT** S.name, S.gpa  
FROM Students S  
WHERE S.dept = 'CS'  
[GROUP BY *<column list>*  
[HAVING *<predicate>*] ]  
[ORDER BY *<column list>*]
- Simplest version is straightforward
  - Produce all tuples in the table that satisfy the predicate
  - Output the expressions in the SELECT list
    - Expression can be a column reference, or an arithmetic expression over column refs

# ORDER BY

- ```
SELECT DISTINCT S.name, S.gpa, S.age*2 AS a2
  FROM Students S
 WHERE S.dept = 'CS'
[GROUP BY <column list>
[HAVING <predicate>] ]
ORDER BY S.gpa, S.name, a2;
```
- ORDER BY clause specifies that output **should be sorted**
  - Lexicographic ordering again!
- Obviously must refer to columns in the output
  - Note the AS clause for naming output columns!

# ORDER BY

- SELECT DISTINCT S.name, S.gpa  
FROM Students S  
WHERE S.dept = 'CS'  
[GROUP BY *<column list>*  
[HAVING *<predicate>*] ]  
**ORDER BY** S.gpa **DESC**, S.name **ASC**;
- Ascending order by default, but can be overridden
  - DESC flag for descending, ASC for ascending
  - Can mix and match, lexicographically

# Aggregates

- ```
SELECT [DISTINCT] AVERAGE(S.gpa)
      FROM Students S
      WHERE S.dept = 'CS'
      [GROUP BY <column list>
      [HAVING <predicate>] ]
      [ORDER BY <column list>]
```
- Before producing output, compute a **summary (a.k.a. an aggregate)** of some arithmetic expression
- Produces **1 row of output**
  - with one column in this case
- Other aggregates: SUM, COUNT, MAX, MIN
- Note: can use DISTINCT inside the agg function
  - SELECT COUNT(DISTINCT S.name) FROM Students S
  - vs. SELECT DISTINCT COUNT (S.name) FROM Students S;

# GROUP BY

- ```
SELECT [DISTINCT] AVERAGE(S.gpa), S.dept
FROM Students S
[WHERE <predicate>]
GROUP BY S.dept
[HAVING <predicate>]
[ORDER BY <column list>]
```
- **Partition** the table into **groups** that have the same value on GROUP BY columns a.k.a. *grouping key*
  - Can group by a list of columns
- Produce **an aggregate** result **per group**
  - Cardinality of output = # of distinct group values
- Note: can put grouping keys in SELECT list
  - For aggregate queries, SELECT list can **contain aggs and grouping keys only!**

# HAVING

- ```
SELECT [DISTINCT] AVERAGE(S.gpa), S.dept
FROM Students S
[WHERE <predicate>]
GROUP BY S.dept
HAVING COUNT(*) > 25
[ORDER BY <column list>]
```
- The HAVING predicate is applied *after* grouping and aggregation
  - Hence can contain anything that could go in the SELECT list
  - i.e. aggs or GROUP BY columns (i.e. grouping keys)
- HAVING can only be used in aggregate queries
- It's an optional clause



# Putting it all together

- ```
SELECT S.dept, COUNT(*)  
  FROM Students S  
 WHERE S.gender = "F"  
 GROUP BY S.dept  
  HAVING COUNT(*) > 25  
 ORDER BY S.dept;
```

# Multi-relation Queries

- Interesting queries often combine data from **more than one relation**.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>” .

# Types of multi-relation queries



- **Join**—a relational operation that causes two or more tables with a **common domain** to be **combined** into a **single table** or **view**

# Types of multi-relation Queries

- Join

- ➔ • **Inner-join**—a join that will only return rows from each table that **have matching rows** in the other
  - For each **customer** who placed an **order**, what is the customer's name and order number?

```
SELECT Customer_T.CustomerID, Order_T.CustomerID,  
       CustomerName, OrderID  
FROM Customer_T INNER JOIN Order_T ON  
       Customer_T.CustomerID = Order_T.CustomerID  
ORDER BY OrderID;
```

```
SELECT Customer_T.CustomerID, Order_T.CustomerID,  
       CustomerName, OrderID  
FROM Customer_T, Order_T  
       WHERE Customer_T.CustomerID = Order_T.CustomerID  
ORDER BY OrderID
```

# Types of multi-relation Queries

- Join

- Inner-join

- ➔ • **Outer join**—a join in which rows that do not have matching values in common columns are **nonetheless** included in the result table
  - List the customer name, ID number, and order number for all customers. Include customer information even for customers that do not have an order?

```
SELECT Customer_T.CustomerID, Order_T.CustomerID,  
       CustomerName, OrderID  
FROM Customer_T, Order_T  
WHERE Customer_T.CustomerID = Order_T.CustomerID  
ORDER BY OrderID
```

# Types of multi-relation Queries


- Join

- Inner-join

- ➔ • **Outer join**—a join in which rows that do not have matching values in common columns are **nonetheless** included in the result table
  - List the customer name, ID number, and order number for all customers. Include customer information even for customers that do not have an order?

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
FROM Customer_T LEFT OUTER JOIN Order_T  
WHERE Customer_T.CustomerID = Order_T.CustomerID;
```


# Outline

- Quick SQL review
- Query processing overview 
- Relational Algebra review
- Introduction to Apache Calcite
- Join processing

# Query Processing Overview

- Query processing requires translating SQL to a special **internal language**
  - Query Plans
- The **query executor** is an **interpreter** for **query plans**
- Think of **query plans** as “box-and-arrow” **dataflow** diagrams
  - Each **box** implements a (*relational algebra*) **operator**
  - **Edges** represent a **flow of tuples** (columns as specified)

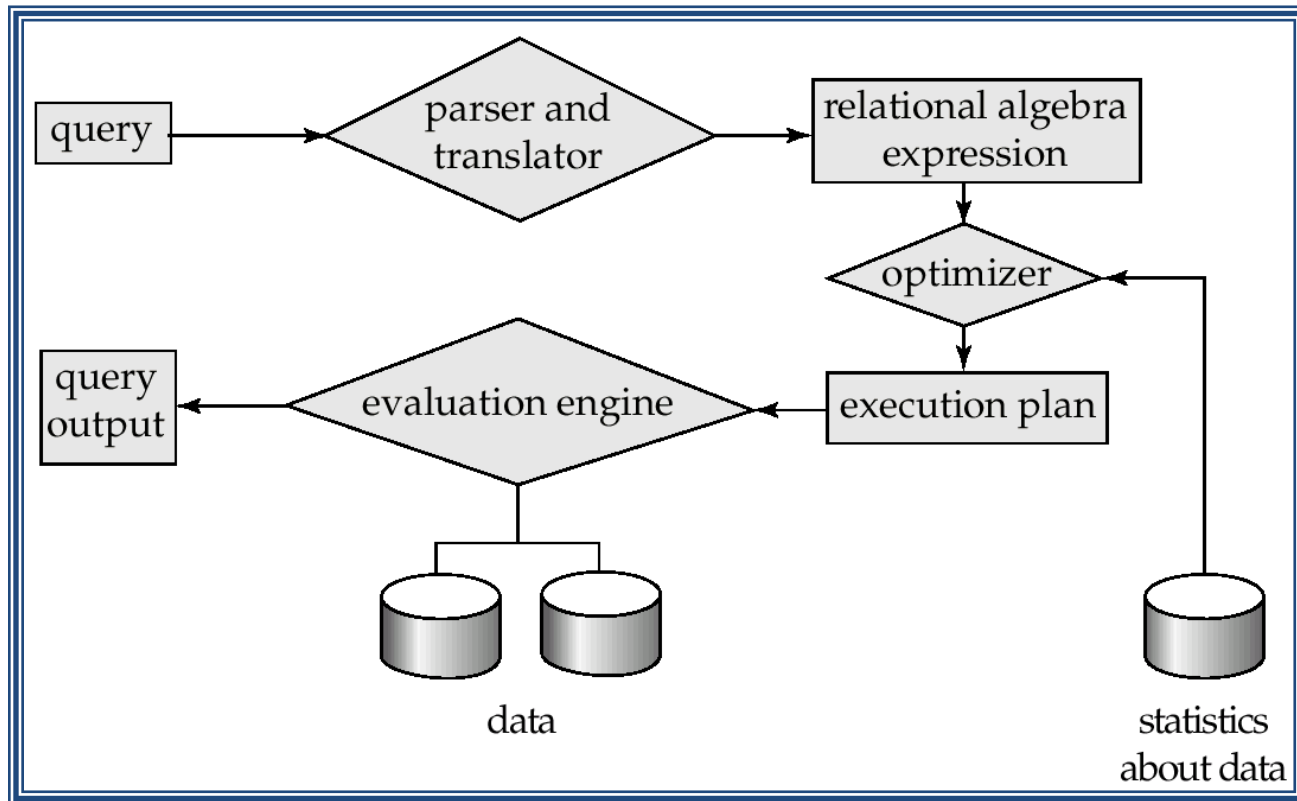
SELECT DISTINCT name, gpa  
FROM Students





# Basic Steps in Query Processing

1. Parsing and translation
2. Optimization
3. Evaluation



# Basic Steps in Query Processing (Cont.)

- Parsing and translation
  - Translate the query into its **internal form**. This is then translated into **relational algebra** expressions (query plans)
  - Parser checks syntax, verifies relations
- Query optimization
  - The query optimizer chooses the **best query plan** (execution plan)
- Execution
  - The query-execution engine takes a execution plan, **executes** that plan, and returns the answers to the query.

# Query plan

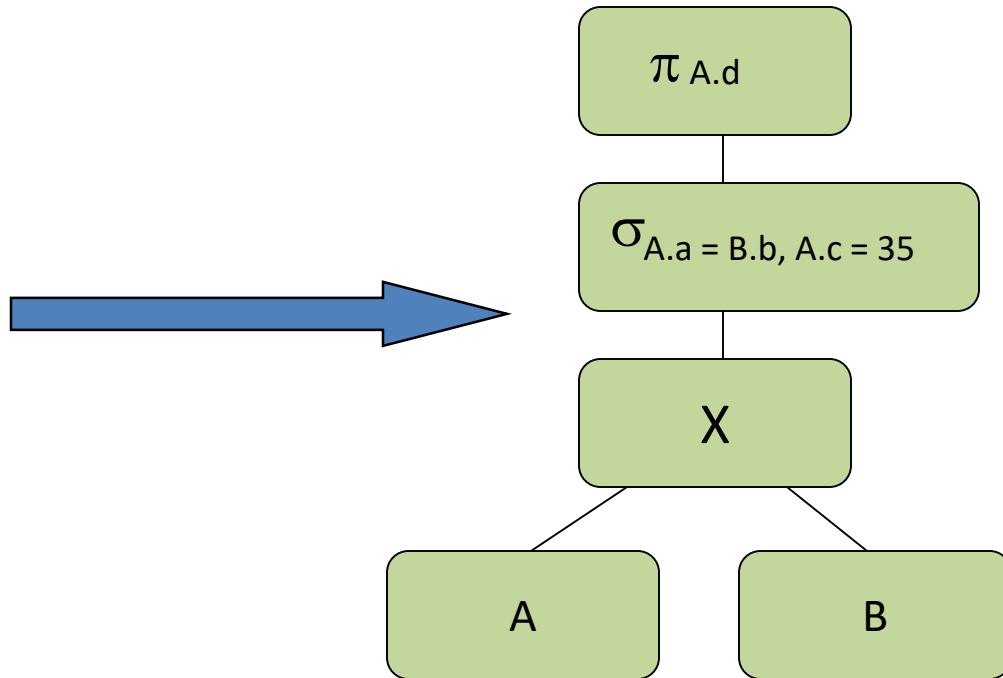
## TABLES:

A(a,c,d)

B(b,e,f)

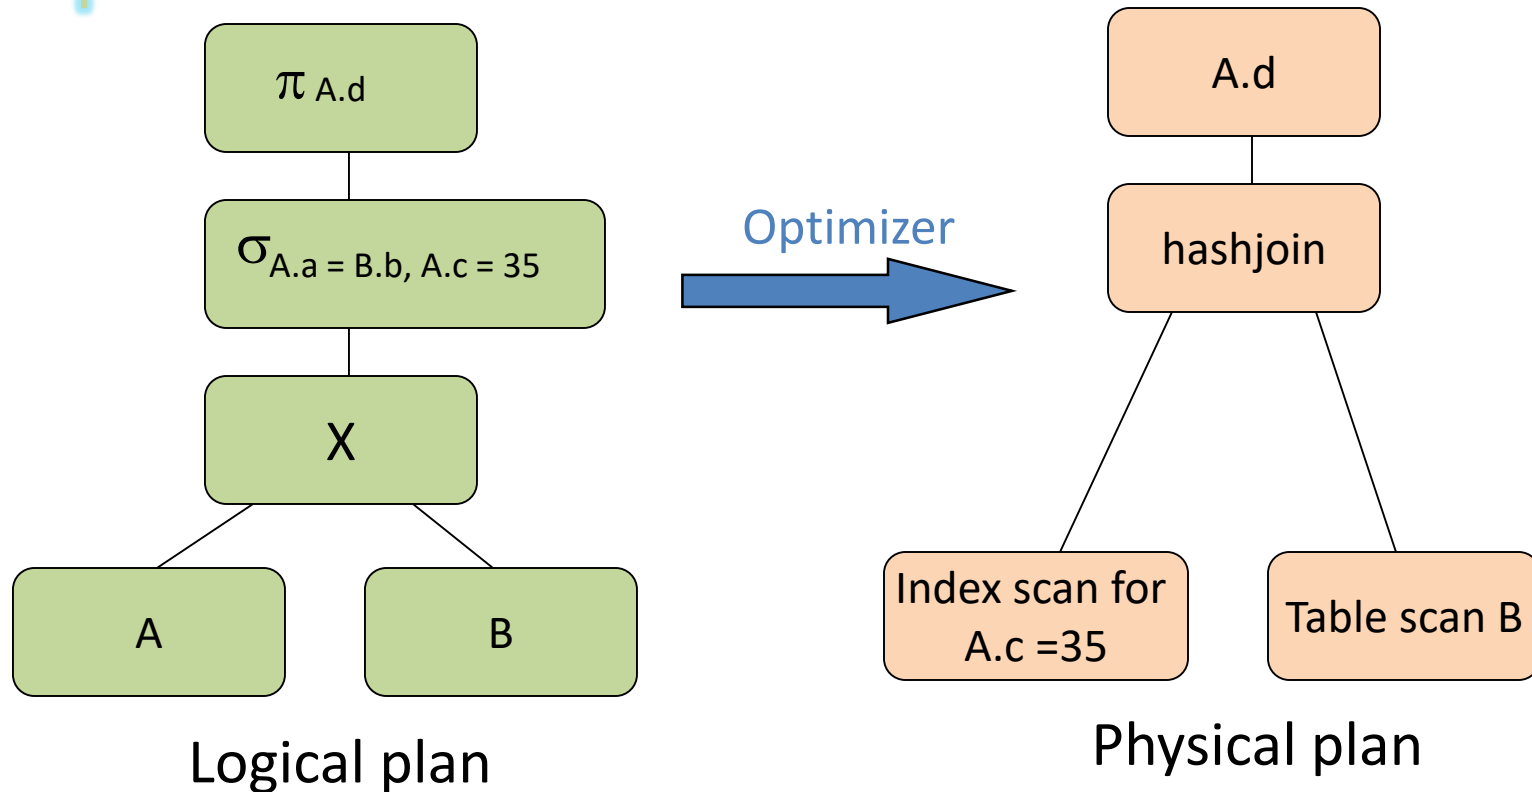
## QUERY:

**SELECT A.d**  
**FROM A, B**  
**WHERE A.a = B.b**  
**AND A.c = 35**



- Relational algebra for SQL very well understood

# Query Optimization





- logical, e.g., push down cheap predicates
- enumerate alternative plans, apply cost model
- use search heuristics to find cheapest plan

# Iterators

- The **relational operators** can be implementation of the interface **Iterator**:

```
interface Iterator {
    void init();
    tuple next();
    void close();
    // additional states go here
    ...
}
```

```
class TableScanIterator implements
Iterator {
    void init() {...}
    tuple next() {...}
    void close() {...}
    // additional states go here
    ...
}
```



UML relations notation

# Outline


- Quick SQL review
- Query processing overview
- Relational Algebra review ←
- Introduction to Apache Calcite
- Join processing

# Relational Algebra summary

- Selection:  $\sigma_p(R)$ 
  - Returns all rows in R that satisfy p
- Projection:  $\pi_C(R)$ 
  - Returns all rows in R projected to columns in C
    - In strict relational model, remove duplicate rows
    - In SQL, preserve duplicates (multiset/bag semantics)
- Cartesian Product:  $R \times S$
- Union:  $R \cup S$       Intersection:  $R \cap S$ 
  - Note: R, S must have matching schema
- Join:  $R \bowtie_p S = \sigma_p(R \times S)$

# Unary operators: select ( $\sigma$ )

- $\sigma_p(R)$  outputs tuples of  $R$  which satisfy  $P$
- same schema as  $R$



*Removes unwanted rows from relation*



# Unary operators: select ( $\sigma$ ) example

## Employees

| Surname | FirstName | Age | Salary |
|---------|-----------|-----|--------|
| Smith   | Mary      | 25  | 2000   |
| Black   | Lucy      | 40  | 3000   |
| Verdi   | Nico      | 36  | 4500   |
| Smith   | Mark      | 40  | 3900   |


Q. Show the employees with salary greater than 4000 and age less than 30

$\sigma_{\text{Age} < 30 \vee \text{Salary} > 4000}$  (**Employees**)

| Surname | FirstName | Age | Salary |
|---------|-----------|-----|--------|
| Smith   | Mary      | 25  | 2000   |
| Verdi   | Nico      | 36  | 4500   |

# Unary operators: project ( $\pi$ )

- $\pi_Y(R)$  outputs a subset  $Y$  of the set of attributes  $X$  of relation  $R$



*Removes unwanted columns from relation*

# Unary operators: project ( $\pi$ ) example

## Employees

| Surname | FirstName | Department | Head     |
|---------|-----------|------------|----------|
| Smith   | Mary      | Sales      | De Rossi |
| Black   | Lucy      | Sales      | De Rossi |
| Verdi   | Mary      | Personnel  | Fox      |
| Smith   | Mark      | Personnel  | Fox      |


Q. Show the surname and firstname of the employees

$\pi_{\text{Surname, FirstName}}(\text{Employees})$

| Surname | FirstName |
|---------|-----------|
| Smith   | Mary      |
| Black   | Lucy      |
| Verdi   | Mary      |
| Smith   | Mark      |

# Additive operators ( $\cup$ , $\cap$ , $-$ )

- Standard set operators
- Operate on tuples within input relations, but not on schema



# Additive operators: Union (U)

## Graduates

| Number | Surname  | Age |
|--------|----------|-----|
| 7274   | Robinson | 37  |
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |

## Managers

| Number | Surname  | Age |
|--------|----------|-----|
| 9297   | O'Malley | 56  |
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |

## Graduates $\cup$ Managers

| Number | Surname  | Age |
|--------|----------|-----|
| 7274   | Robinson | 37  |
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |
| 9297   | O'Malley | 56  |

# Additive operators: Intersection ( $\cap$ )

## Graduates

| Number | Surname  | Age |
|--------|----------|-----|
| 7274   | Robinson | 37  |
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |

## Managers

| Number | Surname  | Age |
|--------|----------|-----|
| 9297   | O'Malley | 56  |
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |

## Graduates $\cap$ Managers

| Number | Surname  | Age |
|--------|----------|-----|
| 7432   | O'Malley | 39  |
| 9824   | Darkes   | 38  |

# Cartesian product ( $\times$ )

- The outcome of combining every record in R with every record in S
- $T = R \times S$  contains every pairwise combination of R and S tuples
  - $\text{schema}(T) = \text{schema}(R) \cup \text{schema}(S)$

|   |  |
|---|--|
| 1 |  |
| 2 |  |
| 3 |  |

|  |   |
|--|---|
|  | A |
|  | B |
|  | C |

$\xrightarrow{\times}$

|   |  |  |   |
|---|--|--|---|
| 1 |  |  | A |
| 1 |  |  | B |
| 1 |  |  | C |
| 2 |  |  | A |
| 2 |  |  | B |
| 2 |  |  | C |
| 3 |  |  | A |
| 3 |  |  | B |
| 3 |  |  | C |

# Cartesian product ( $\times$ ) example

## Employees

| Employee | Project |
|----------|---------|
| Smith    | A       |
| Black    | A       |
| Black    | B       |

## Projects

| Code | Name  |
|------|-------|
| A    | Venus |
| B    | Mars  |

## Employees $\times$ Projects

| Employee | Project | Code | Name  |
|----------|---------|------|-------|
| Smith    | A       | A    | Venus |
| Black    | A       | A    | Venus |
| Black    | B       | A    | Venus |
| Smith    | A       | B    | Mars  |
| Black    | A       | B    | Mars  |
| Black    | B       | B    | Mars  |




# Natural join ( $\bowtie$ )

- $T = R \bowtie S$  merges tuples from R and S having equal values where their schemas overlap (**join attributes**)
  - T Schema: Union of schemas  
 $schema(R) \cap schema(S) \neq \emptyset$

- Special cases

- No schema overlap:  $\times$
- Full schema overlap:  $\cap$



# Natural join ( $\bowtie$ ) example

$r_1$

| <u>Employee</u> | Department |
|-----------------|------------|
| Smith           | sales      |
| Black           | production |
| White           | production |

$r_2$

| <u>Department</u> | Head  |
|-------------------|-------|
| production        | Mori  |
| sales             | Brown |

$r_1 \bowtie r_2$

| <u>Employee</u> | <u>Department</u> | <u>Head</u> |
|-----------------|-------------------|-------------|
| Smith           | sales             | Brown       |
| Black           | production        | Mori        |
| White           | production        | Mori        |

# Theta join

- Written as  $T = R \bowtie_{\theta} S$ 
  - Outputs pairwise combinations of tuples which satisfy  $\theta$
- Most general join
  - Arbitrary join **predicate** (not just equality)

# Theta join example

**Car**

| Car  | CarPrice |
|------|----------|
| CarA | 20000    |
| CarB | 30000    |
| CarC | 50000    |

**Boat**

| Boat  | BoatPrice |
|-------|-----------|
| BoatA | 10000     |
| BoatB | 40000     |
| BoatC | 60000     |

Q. select the cars and boats where car price is higher than boat price

**Car** ⋈<sub>CarPrice > BoatPrice</sub> **Boat**

| Car  | CarPrice | Boat  | BoatPrice |
|------|----------|-------|-----------|
| CarA | 20000    | BoatA | 10000     |
| CarB | 30000    | BoatA | 10000     |
| CarC | 50000    | BoatA | 10000     |
| CarC | 50000    | BoatB | 40000     |

# Equijoin

- Special case of theta join
- Written as  $R \bowtie_{A=X, B=Y, \dots} S$ 
  - Attribute names in R and S can differ
  - Still compare values for equality
- Like natural join, but using arbitrary attributes
  - Very common due to *foreign keys* in relations

# Equijoin example

## Employees

| Employee | Project |
|----------|---------|
| Smith    | A       |
| Black    | A       |
| Black    | B       |

## Projects

| Code | Name  |
|------|-------|
| A    | Venus |
| B    | Mars  |


Q. select employees and the projects they work on

## Employees ⋈<sub>Project=Code</sub> Projects

| Employee | Project | Code | Name  |
|----------|---------|------|-------|
| Smith    | A       | A    | Venus |
| Black    | A       | A    | Venus |
| Black    | B       | B    | Mars  |

# Outer join ( $\bowtie$ )

- $T = R \bowtie S$  computes the “outer” join of R (left) and S (right)
  - Like normal join, but all tuples from R and S appear in output
  - Pad (left, right, or all) dangling tuples with  $\perp$  or **NULL**
    - **LEFT** — Tuples in inner join padded with tuples in R that have no matching tuples in S.
    - **RIGHT** — Tuples in inner join padded with tuples in S that have no matching tuples in R.
    - **FULL** — Tuples in inner join padded with tuples in R that have no matching tuples in S and tuples in S that have no matching tuples in R.
  - *Natural, equi-, and theta- variants still apply*
  - $|T| \geq \max(|R|, |S|)$



# Outer join ( $\bowtie$ ) examples

$r_1$

| Employee | Department |
|----------|------------|
| Smith    | sales      |
| Black    | production |
| White    | production |

$r_2$

| Department | Head  |
|------------|-------|
| production | Mori  |
| purchasing | Brown |

$r_1 \bowtie r_2$

| Employee | Department | Head |
|----------|------------|------|
| Smith    | sales      | NULL |
| Black    | production | Mori |
| White    | production | Mori |

$r_1 \bowtie r_2$

| Employee | Department | Head  |
|----------|------------|-------|
| Black    | production | Mori  |
| White    | production | Mori  |
| NULL     | purchasing | Brown |

$r_1 \bowtie r_2$


| Employee | Department | Head  |
|----------|------------|-------|
| Smith    | Sales      | NULL  |
| Black    | production | Mori  |
| White    | production | Mori  |
| NULL     | purchasing | Brown |



# Sorting ( $\tau$ )

- $\tau_L(R)$  sorts tuples in R on list of attributes L
  - If L is A1, A2, ..., An tuples sorted first by A1. Ties are broken based on A2;...; Ties that remain after An broken arbitrarily.
  - Default: ascending order; With '-' in front: descending order
- Example:  $\tau_{\text{-Count, Make}}(R)$

| Make   | Count |
|--------|-------|
| Toyota | 2     |
| Honda  | 3     |
| Ford   | 3     |



| Make   | Count |
|--------|-------|
| Ford   | 3     |
| Honda  | 3     |
| Toyota | 2     |

Descending count

Alphabetical order  
when count is equal

# Grouping ( $\Gamma$ )

- Aggregate functions
  - min, max, sum, count, average, ...
- $\Gamma_{A,B,C,f(X),g(Y),h(Z)}(R)$  computes aggregate values using some attributes as a grouping key
  - Implicit projection (drops unreferenced attributes)
  - A, B, C is the *grouping key*
  - X, Y, Z are *attributes* to aggregate
  - *f, g, h* are *aggregating* functions to apply

# Grouping Example

- Answer the following query
  - “List employees and their total sales in descending order”
  - $\tau_{\text{Total}}(\rho_{\text{Name,Total}}(\Gamma_{\text{Name,sum(Value)}}(\text{Emp} \bowtie \text{Sales})))$

**Emp**

| EID | Name     |
|-----|----------|
| 1   | Mary     |
| 2   | Xiao     |
| 3   | Jaspreet |

**Sales**

| EID | Value | ... |
|-----|-------|-----|
| 1   | 20    | ... |
| 3   | 10    | ... |
| 3   | 15    | ... |

=

| Name     | Total |
|----------|-------|
| Jaspreet | 25    |
| Mary     | 20    |

# Schema for Additional Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Boats (*bid*: integer, *bname*: string, *color*: string)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)



# Example 1

Find names of sailors who've reserved boat #103


## Example 2

Find sailors who've reserved a red or a green boat

# Exercise

Find sailors who've reserved a red and a green boat

# Outline

- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite   
(Some slides are from J. Halterman)
- Join processing



# What is Apache Calcite?

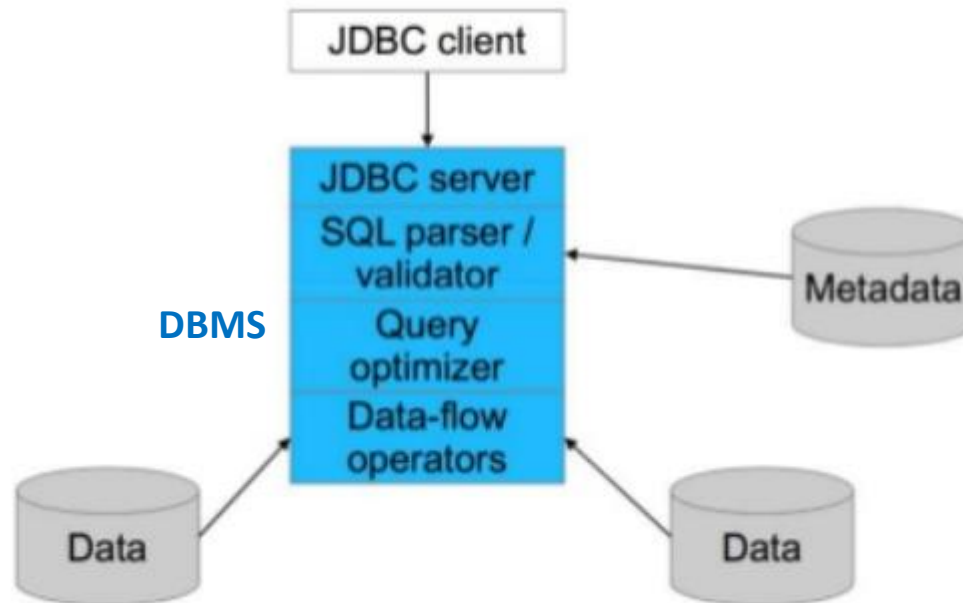
- A framework for building SQL databases
- Developed over more than ten years
- Written in Java
- Previously known as Optiq
- Previously known as Farrago
- Became an Apache project in 2013
- Led by Julian Hyde at Hortonworks

# Projects using Calcite

- Apache Hive
- Apache Drill
- Apache Flink
- Apache Phoenix
- Apache Samza
- Apache Storm
- Apache everything...

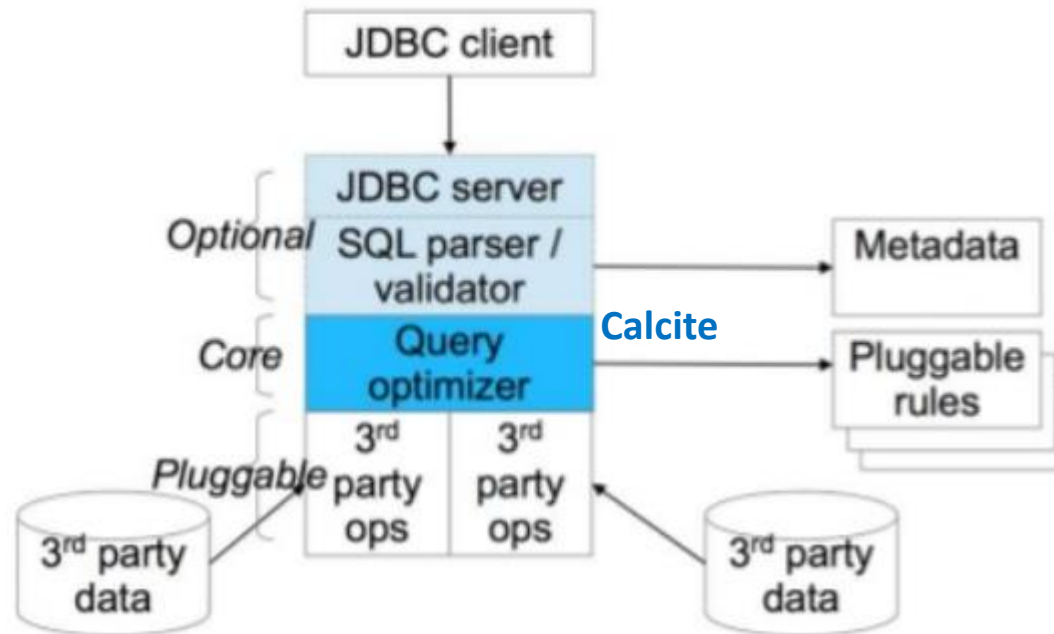
# What is Apache Calcite?

## Conventional DBMS architecture



# What is Apache Calcite?

## Apache Calcite architecture



# Stages of query execution

**01**

Parse

Queries are parsed using a JavaCC generated parser

**02**

Validate

Queries are validated against known database metadata

**03**


Optimize

Logical plans are optimized and converted into physical expressions

**04**

Execute

Physical plans are converted into application-specific executions

- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite
  - Components 
- Join processing

# Components of Calcite

- **Catalog** - Defines metadata and namespaces that can be accessed in SQL queries
- **SQL parser** - Parses valid SQL queries into an abstract syntax tree (AST)
- **SQL validator** - Validates abstract syntax trees against metadata provided by the catalog
- **Query optimizer** - Converts AST into logical plans, optimizes logical plans, and converts logical expressions into physical plans

# Outline

- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite
  - Components
  - Query plan and optimization ←
- Join processing



# Query Plans

- Query plans represent the steps necessary to execute a query

# Query Plans

- Query plans represent the steps necessary to execute a query

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id  
FROM users u INNER JOIN orders o ON u.id = o.user_id  
WHERE u.id > 50
```

# Query Plans

- Query plans represent the steps necessary to execute a query

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id  
FROM users u INNER JOIN orders o ON u.id = o.user_id  
WHERE u.id > 50
```

table scan →

↑ table scan

# Query Plans

- Query plans represent the steps necessary to execute a query

```

SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id
FROM users u INNER JOIN orders o ON u.id = o.user_id
WHERE u.id > 50
  
```

table scan →  
 inner join  
 table scan

# Query Plans

- Query plans represent the steps necessary to execute a query

```


SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id
FROM users u INNER JOIN orders o ON u.id = o.user_id
WHERE u.id > 50
  
```

Diagram illustrating the query plan components:

- table scan**: Points to the `FROM users u` clause.
- filter**: Points to the `WHERE u.id > 50` clause.
- inner join**: Points to the `INNER JOIN orders o` clause.
- table scan**: Points to the `orders o` table in the join clause.

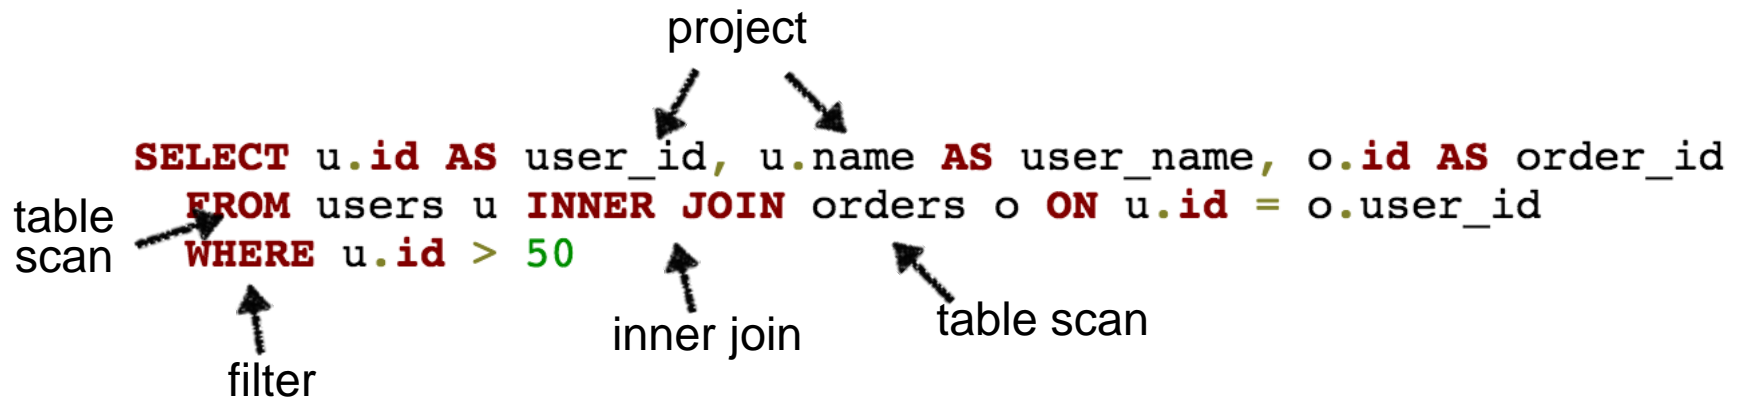
# Query Plans

- Query plans represent the steps necessary to execute a query



# Query Plans

- Query plans represent the steps necessary to execute a query



```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])
  LogicalFilter(condition=[>($0, 50)])
    LogicalJoin(condition=[=($0, $6)], joinType=[inner])
      LogicalTableScan(table=[[USERS]])
      LogicalTableScan(table=[[ORDERS]])
```

# Query Optimization

- Optimize logical plan
- Goal is typically to try to reduce the amount of data that must be processed early in the plan
- Convert logical plan into a physical plan
- Physical plan is engine specific and represents the physical execution stages



# Query Optimization

- Prune unused fields
- Merge projections
- Convert subqueries to joins
- Reorder joins
- Push down projections
- Push down filters

# Query Optimization

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id  
FROM users u INNER JOIN orders o ON u.id = o.user_id  
WHERE u.id > 50
```

# Query Optimization

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id  
FROM users u INNER JOIN orders o ON u.id = o.user_id  
WHERE u.id > 50
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])  
  LogicalFilter(condition=[>([0], 50)])  
    LogicalJoin(condition=[=[0, 6]], joinType=[inner])  
      LogicalTableScan(table=[[USERS]])  
      LogicalTableScan(table=[[ORDERS]])
```

# Query Optimization

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id  
FROM users u INNER JOIN orders o ON u.id = o.user_id  
WHERE u.id > 50
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])  
  LogicalFilter(condition=[>(0, 50)])  
    LogicalJoin(condition=[=(0, 6)], joinType=[inner])  
      LogicalTableScan(table=[[USERS]])  
        LogicalTableScan(table=[[ORDERS]])
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])  
  LogicalJoin(condition=[=(0, 6)], joinType=[inner])  
    LogicalProject(ID=[0], NAME=[1])  
      LogicalFilter(condition=[>(0, 50)])  
        LogicalTableScan(table=[[USERS]])  
    LogicalProject(ID=[0], USER_ID=[1])  
      LogicalTableScan(table=[[ORDERS]])
```

# Query Optimization

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id
FROM users u INNER JOIN orders o ON u.id = o.user_id
WHERE u.id > 50
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])
  LogicalFilter(condition=[>(0, 50)])
    LogicalJoin(condition=[=(0, 6)], joinType=[inner])
      LogicalTableScan(table=[[USERS]])
      LogicalTableScan(table=[[ORDERS]])
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])
  LogicalJoin(condition=[=(0, 6)], joinType=[inner])
    LogicalProject(ID=[0], NAME=[1])
      LogicalFilter(condition=[>(0, 50)])
        LogicalTableScan(table=[[USERS]])
    LogicalProject(ID=[0], USER_ID=[1])
      LogicalTableScan(table=[[ORDERS]])
```

push down  
project

# Query Optimization

```
SELECT u.id AS user_id, u.name AS user_name, o.id AS order_id
FROM users u INNER JOIN orders o ON u.id = o.user_id
WHERE u.id > 50
```


```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])
  LogicalFilter(condition=[>(0, 50)])
    LogicalJoin(condition=[=(0, 6)], joinType=[inner])
      LogicalTableScan(table=[[USERS]])
      LogicalTableScan(table=[[ORDERS]])
```

```
LogicalProject(user_id=[0], user_name=[1], order_id=[5])
  LogicalJoin(condition=[=(0, 6)], joinType=[inner])
    LogicalProject(ID=[0], NAME=[1])
      LogicalFilter(condition=[>(0, 50)])
        LogicalTableScan(table=[[USERS]])
    LogicalProject(ID=[0], USER_ID=[1])
      LogicalTableScan(table=[[ORDERS]])
```

push down  
project

push down  
filter

# Outline

- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite
  - Components
  - Query plan and optimization
  - Key concepts 
- Join processing

# Key Concepts

- **Relational algebra** - `RelNode`
  - A relational algebra expression
- **Row expressions** - `RexNode`
  - A row-level expression (e.g. Projection fields, Filter condition)
- **Traits** - `RelTrait`
  - A trait of a relational expression that does not alter execution
- **Conventions** –
  - used to represent a single data source (`SparkConvention`, `JdbcConvention` etc)
- **Rules** - `RelOptRule`
  - Rules are used to modify query plans
- **Planners** - `RelOptPlanner`
  - Represents the query planner



# Key Concepts

Relational algebra

`RelNode`

Row expressions

`RexNode`

Traits

`RelTrait`

Conventions

`Convention`

Rules

`RelOptRule`


Planners

`RelOptPlanner`

# Relational Algebra

- `RelNode` represents a relational expression
- Largely equivalent to Spark's `DataFrame` methods
- Logical algebra
- Physical algebra

# Outline

- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite
  - Components
  - Query plan and optimization
  - Key concepts
  - Relational algebra builder 
- Join processing

# Relational Algebra APIs

| Operation                                                                | RA Operator                                    | Relbuilder method | Method argument                                                                |
|--------------------------------------------------------------------------|------------------------------------------------|-------------------|--------------------------------------------------------------------------------|
| Table scan                                                               |                                                | scan              |                                                                                |
| Project                                                                  | $\pi_c(R)$                                     | project           |                                                                                |
| Select                                                                   | $\sigma_p(R)$                                  | filter            |                                                                                |
| Join (inner)<br>Join(left out.)<br>Join (right out.)<br>Join (full out.) | $\bowtie$<br>$\ltimes$<br>$\rtimes$<br>$\Join$ | join              | JoinRelType.INNER<br>JoinRelType.LEFT<br>JoinRelType.RIGHT<br>JoinRelType.FULL |
| Union                                                                    | $\cup$                                         | union             |                                                                                |
| Intersection                                                             | $\cap$                                         | intersect         |                                                                                |
| Grouping                                                                 | $\Gamma$                                       | aggregate         |                                                                                |
| Sorting                                                                  | $\tau$                                         | sort              |                                                                                |
| Rename                                                                   | $\rho$                                         | as                |                                                                                |
| And                                                                      | $\wedge$                                       | and               |                                                                                |
| Or                                                                       | $\vee$                                         | or                |                                                                                |

# Steps to using building and running Relational Algebra expressions

```
import org.apache.calcite.tools.RelBuilder;
import org.apache.calcite.tools.RelRunners;

// Create a builder. The config contains a schema mapped
final FrameworkConfig config = buildConfig(); // see code
final RelBuilder builder = RelBuilder.create(config);

// Build RA expression for query: select * from COURSE
builder.scan("COURSE");

// Returns the final relational algebra expression
final RelNode node = builder.build();

// execute the query plan
try {
    final PreparedStatement preparedStatement =
        RelRunners.run(node, calConn);

    ResultSet rs = preparedStatement.executeQuery();
    while (rs.next()) { // do something }
    rs.close();
} catch (SQLException e) {
}
```

# Examples of building RA expression: example 1

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show the title of the course where courseid = 2
select TITLE from COURSE where COURSEID = 2
```

```
// Build RA expression for the above query
builder
  .scan("COURSE")
  .filter( builder.equals(builder.field("COURSEID"), builder.literal(2)) )
// or
//.filter( builder.call(SqlStdOperatorTable.EQUALS,
builder.field("COURSEID"), builder.literal(2) ) )
  .project(builder.field("TITLE"));

// See Javadoc for RelBuilder APIs
// https://calcite.apache.org/apidocs/org/apache/calcite/tools/RelBuilder.html
```

# Examples of building RA expression: example 2

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)  
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show the coursed and title of the first 5 courses sorted by the  
course
```

```
// Build RA expression for the above query
```

```
builder
```

```
.scan("COURSE")
```

```
.sort(  builder.field("COURSEID")  )
```

```
.limit(0, 5) // offset 0, limit 5
```

```
.project(builder.field("COURSEID"), builder.field("TITLE"));
```

```
// See Javadoc for RelBuilder APIs
```

```
// https://calcite.apache.org/apidocs/org/apache/calcite/tools/RelBuilder.html
```

# Examples of building RA expression: example 3

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show the number of courses in each course category
SELECT CATEGORYID, count(*) AS C,
FROM COURSE
GROUP BY CATID
```

```
// Build RA expression for the above query
```

```
builder
  .scan("COURSE")
  .aggregate(builder.groupKey("CATEGORYID"),
              builder.count(false, "C", builder.field("COURSEID"))
  );
```



# Examples of building RA expression: example 3

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)  
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show the number of courses in each course category where number of  
courses is greater than 1
```

```
SELECT CATEGORYID, count(*) AS C,  
FROM COURSE  
GROUP BY CATID  
HAVING C > 1
```

```
// Build RA expression for the above query
```

```
builder  
  .scan("COURSE")  
  .aggregate(builder.groupKey("CATEGORYID"),  
             builder.count(false, "C", builder.field("COURSEID")) )  
  )  
  .filter( builder.call(SqlStdOperatorTable.GREATER_THAN,  
                      builder.field("C"), builder.literal(1))) ;
```

# Examples of building RA expression: example 4

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show the title of each course along with the name of the category
SELECT TITLE, CATNAME
FROM COURSE c, CCATEGORY g
WHERE c.CATEGORYID = g.CATID
```

```
// Build RA expression for the above query
```

```
builder
```

```
.scan("COURSE").as("c")
```

```
.scan("CCATEGORY").as("g")
```

```
.join(JoinRelType.INNER)
```

```
.filter( builder.equals(builder.field("c", "CATEGORYID"),
builder.field("g", "CATID")))
```

```
// Syntax:.filter (predicate1, predicate2); where "," implies AND
```

```
.project(builder.field("TITLE"), builder.field("CATNAME"));
```

# Examples of building RA expression: example 5

## Schema

```
COURSE(COURSEID:int, TITLE:string, CATEGORYID:int)
CCATEGORY(CATID:int, CATNAME:string)
```

## Query

```
-- Show all categories from COURSE and CCATEGORY
SELECT CATEGORYID FROM COURSE
Union
SELECT CATID from CCATEGORY
```

```
// Build RA expression for the above query
```

```
builder
.scan("COURSE").project(builder.field("CATEGORYID"))
.scan("CCATEGORY").project(builder.field("CATID"))
.union(true, 1);
```

# Outline

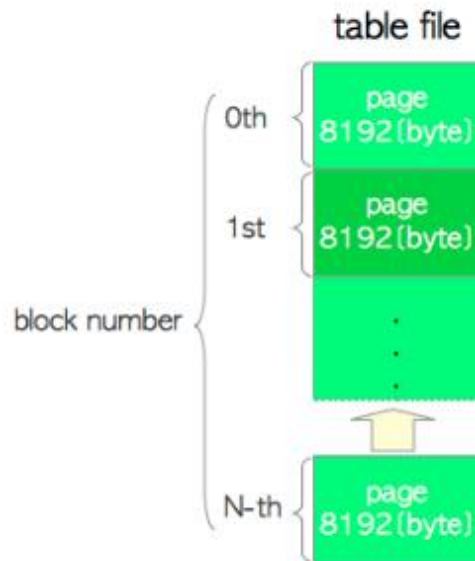
- Quick SQL review
- Query processing overview
- Relational Algebra review
- Introduction to Apache Calcite
- Join processing ←

# Joins - multi-table queries

- Joins are very common
- Joins are **very expensive** (worst case: cross product!)
- Many approaches to reduce join cost
  - (Block) nested loops
  - Indexed nested loops
  - Sort/Merge Join
  - Hash Join
- Main goal: minimize I/O cost

# ASIDE: Internal layout of a table file

- Inside the data file (heap table and index), it is divided into **pages** (or **blocks**) of fixed length, default is **8192 byte** (PostgreSQL)
- Those pages within each file are numbered sequentially from 0, and such numbers are called as **block numbers**.
- If the file has been filled up, database adds a **new empty page** to the **end of the file** to increase the file size.



# Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)  
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

# Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.



# Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- In algebra:  $R \bowtie S$ . Common! Must be carefully optimized.  $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.
- Assume:
  - $M$  pages in  $R$ ,  $p_R$  tuples per page
  - $N$  pages in  $S$ ,  $p_S$  tuples per page.
  - In our examples,  $R$  is Reserves and  $S$  is Sailors.
- *Cost metric* : # of I/Os. We will ignore output costs.

# Nested Loop Join (NLJ)

Basic Join Algorithm

*Input: Relations R and S*

*Output: Joined relation T*


T=empty set

FOR EACH tuple r in R

FOR EACH tuple s in S

IF  $r.a == s.a$

THEN append  $r || s$  to T



# Nested Loops Join (NLJ)

```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
- How much does this Cost?

$$(p_R * M) * N + M = 100 * 1000 * 500 + 1000 = 50,001,000 \text{ I/Os.}$$

– At 10ms/IO, Total: ???

# Index Nested Loop Join (INLJ)

Basic Join Algorithm

*Input: Relations R and S*

*Output: Joined relation T*

T=empty set


FOR EACH tuple r in R

FOR EACH tuple s returned

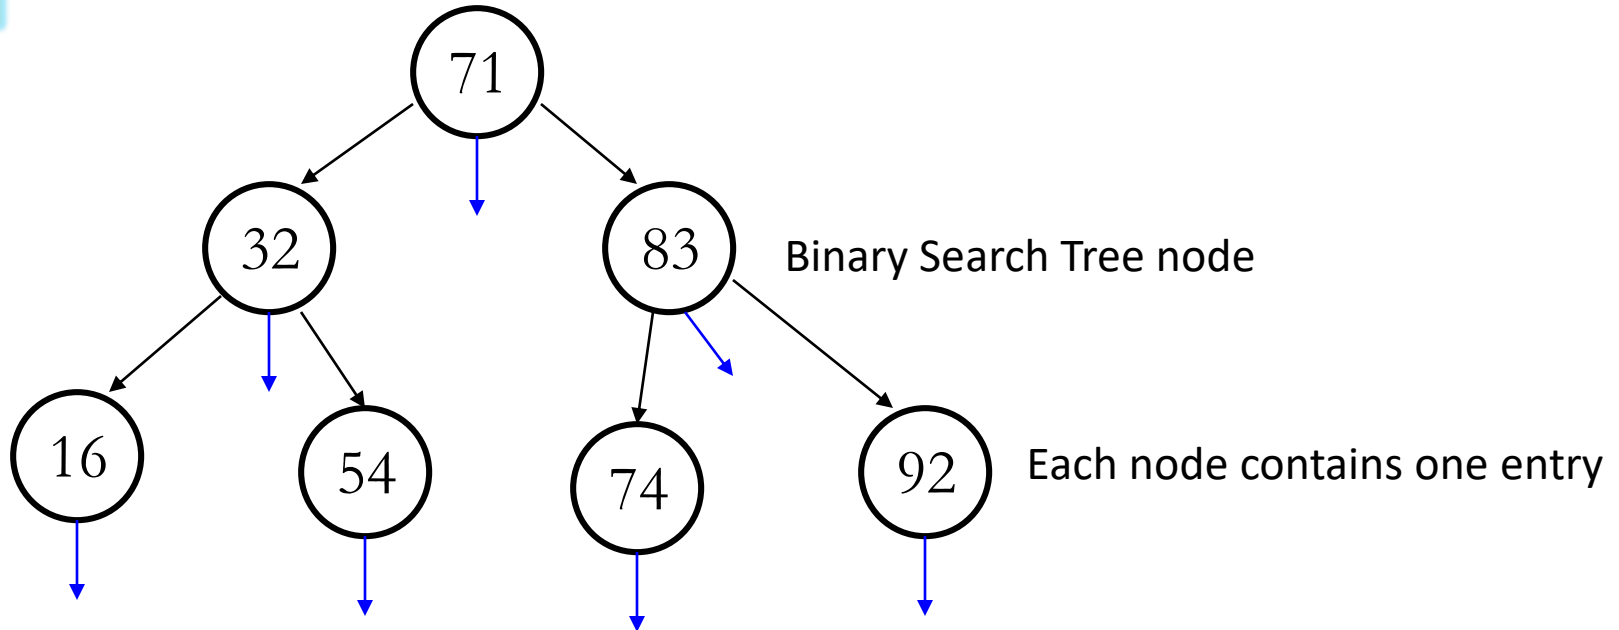
by  $S_{\text{index}}$  where  $r.a == s.a$

Append  $r || s$  to T

Assumption: an **index** exists  
on the join column S.a




# Quick review: Persistent Binary Search Tree?




# Quick review: B<sup>+</sup> Tree Index

- Minimum 50% occupancy (except for root). Each node contains between  $t$  and  $2t$  entries. The parameter  $t$  is called the *order* or *minimum degree* of the tree.
- Height of a B<sup>+</sup>-tree:  $h = \log_t N$   
( $N$  = total number of keys)



# Quick review: B<sup>+</sup> Tree Index

- To **amortize** disk cost, data is transferred in **large chunks**.
  - Think of B-Tree as a **BST with very fat nodes**
  - Each node can store keys for about 50-2000 items, and will have the similar branching factor



- With a branching factor of 1001 (1000 keys per node), 1 billion keys can be accessed by a tree of height 2.
  - Just 2 disk accesses!**

# Index Nested Loops Join (INLJ)

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

- If there is an **index** on the join column of one relation (say S), can make it the inner and exploit the index.
  - Cost:  $M + (M * p_R) * \text{cost of finding matching S tuples}$
- For each tuple, cost of probing index is about **2 - 4 I/Os** for B<sup>+</sup> tree.
- For each Reserves tuple: 2 I/Os to get data entry in index, plus 1 I/O to get (the exactly one) matching Sailors tuple.
  - Cost =  $1000 + (100 * 1000) * 3 = 301,000$  I/Os.
  - Cost (if, **index fits in memory**) =  $1000 + (100 * 1000) = 101,000$  I/Os



# Other Join Approaches


- Sort-merge join
  - Sort both relations first.
  - Scan through both relations only once.
- Hash-based join
  - No need to sort both relations first.
  - Relations have hash indices.
  - Use hash to find the corresponding records.

# Sort-merge Join

Input: Sorted relations R and S

Output: Joined relation T

- General scheme:
  - Do { Advance scan of R until current R-tuple  $\geq$  current S tuple;  
Advance scan of S until current S-tuple  $\geq$  current R tuple; }  
Until current R tuple = current S tuple.
  - At this point, all R tuples with same value in R and all S tuples with same value in S *match*;  
output  $\langle r, s \rangle$  for *all* pairs of such tuples.
    - Like a mini nested loops
  - Then resume scanning R and S.



# Cost of Sort-Merge Join

|            |        |        |      | <u>sid</u> | <u>bid</u> | <u>day</u> | rname  |
|------------|--------|--------|------|------------|------------|------------|--------|
| <u>sid</u> | sname  | rating | age  | 28         | 103        | 12/4/96    | guppy  |
| 22         | dustin | 7      | 45.0 | 28         | 103        | 11/3/96    | yuppy  |
| 28         | yuppy  | 9      | 35.0 | 31         | 101        | 10/10/96   | dustin |
| 31         | lubber | 8      | 55.5 | 31         | 102        | 10/12/96   | lubber |
| 44         | guppy  | 5      | 35.0 | 31         | 101        | 10/11/96   | lubber |
| 58         | rusty  | 10     | 35.0 | 58         | 103        | 11/12/96   | dustin |

- Cost:  $M \log M + N \log N + (M+N)$

– The cost of scanning,  $M+N$

Cost  $\approx 15,950$

# Hash-based Join

*Input: Relations R and S*

*Output: Joined relation T*

**Build** hashtable-R on relation R

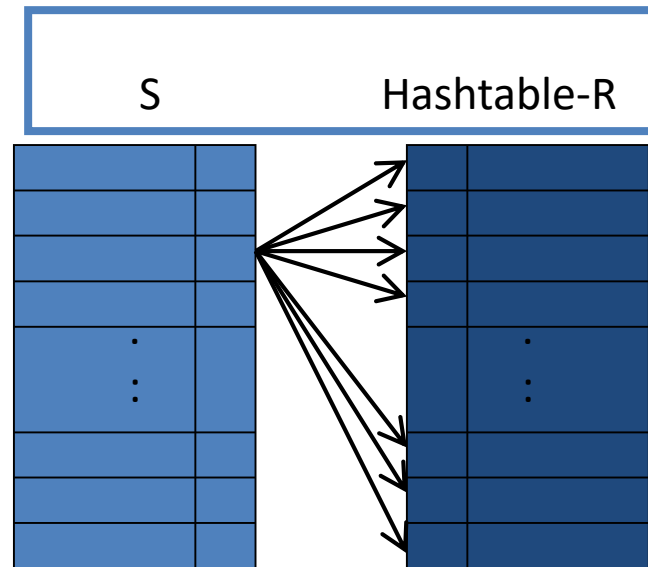
T=empty set

FOR EACH record s in S

    Check hashtable-R,

        IF a match found [i.e.  $h(r.a) == h(s.a)$ ]


        THEN append  $r || s$  to T



- Build in-memory hash table on R
  - R is called the **build** relation of the hash join
- S is called the **probe** relation of the hash join

# Disk-based Hash Join – partitioning phase


- Relations R, S
- Use B buckets
- Read R, hash using *h1*, + write buckets



→ Same for S

# Disk-based Hash Join – join phase

- Read **one** R bucket; build **in-memory** hash table using **h2**  
[R is called the **build** relation of the hash join]
- Read corresponding S bucket + hash probe  
[S is called the **probe** relation of the hash join]



Then repeat for all buckets

# Cost of Hash-Join

- In partitioning phase, read+write both relns;  
 $2(M+N)$ .
- In Join phase, read both relns;  $M+N$  I/Os.
- In our running example, this is a total of 4500 I/Os.

Partitioning:            Read R + write

                              Read S + write

Join:                      Read R, S

Total cost =  $3 \times [1000+500] = 4500$

# Summary of best costs for join algorithms

| Join                        | I/O Cost*            | Time (random access) | Time (sequential access) |
|-----------------------------|----------------------|----------------------|--------------------------|
| Nested loop join (NLJ)      | 50,001,000           | 11.57 days!          | 13.9 hours               |
| Index Nested loop join INLJ | 301,000 or 101,000** | 33.7 minutes         | 1.7 minutes              |
| Sort-merge join             | 15,950               | 5.3 minutes          | 16 seconds               |
| Hash join                   | 4500                 | 1.5 minutes          | 4.5 seconds              |

\* Note: these versions do not use additional in-memory buffer to cache some of the pages. With in-memory buffering, the I/O cost is be significantly improved.

\*\* If the index fits in memory



# Summary of join algorithms

- **NLJ** ok for “small” relations  
(when data fits in memory)
- For **equi-join**, where relations not sorted and no indexes exist,  
**Hash Join** usually best

# Summary of join algorithms

- **Sort-Merge Join** good for non-equi-join (e.g.,  $R1.C > R2.C$ )
- If relations already sorted, use **Merge Join**
- If index exists, it could be useful
  - Depends on expected result size and index clustering
- Join techniques apply to Union, Intersection, Difference

# Summary

- A virtue of relational DBMSs: *queries are composed of a few basic operators*; the implementation of these operators can be carefully tuned (and it is important to do this!).
- Many **alternative implementation techniques** for each operator; no universally superior technique for most operators.
- **Must consider available alternatives for each operation in a query** and **choose best one based** on system statistics, etc. This is part of the broader task of optimizing a query composed of several ops.

# Next

- Parallel Databases..