

# Big Data Systems (CS4545/CS6545) Winter 2021

## Declarative interface on top of batch processing: HiveQL

Suprio Ray


University of New Brunswick, Fredericton

# Acknowledgement

---

Thanks to Tom White (textbook), Raghav Ayyamani and L. Tang for some of the materials in these slides. Also thanks to various articles and research papers.

# Outline

- Hive 
  - Data model
  - System architecture
  - HiveQL
  - File formats
- File-based Data Structures
  - Row-oriented storage formats
  - Column-oriented storage formats

# Why Another Data Warehousing System?


- Problem : Data, data and more data
  - Several TBs of data everyday
- The Hadoop Experience:
  - Uses Hadoop File System (HDFS)
  - Scalable/Available



Src: datanami.com

# Limitations of MapReduce

- Lacks expressiveness
- Map-Reduce can be somewhat hard to program
- For complex jobs:
  - Multiple stage of Map/Reduce functions
  - Just like asking developer to write specific physical execution plan in the database
- Not Reusable
- Error prone



# SQL on MapReduce?


- **Advantages of SQL**
  - SQL has a huge user base
  - SQL is easy to code
- **Solution: Combine SQL and Map-Reduce**
  - Hive on top of Hadoop (open source)



# What is HIVE?

- A system for **managing** and **querying** **unstructured** data **as if** it were **structured**
  - Uses Map-Reduce for execution
  - HDFS for Storage
- Key Building Principles
  - SQL as a **familiar** data warehousing tool
  - **Extensibility** (Pluggable map/reduce scripts in the language of your choice, Rich and **User Defined Data Types**, **User Defined Functions**)
  - **Interoperability** (Extensible Framework to support different file and data formats)
  - **Performance**

# Outline

- Hive
  - Data model 
  - System architecture
  - HiveQL
- File-based Data Structures
  - Row-oriented storage formats
  - Column-oriented storage formats



# Data Model - Tables

- Tables
  - Analogous to **tables** in relational DBs.
  - Each table has a corresponding **directory** in HDFS.
  - Creating a table

```
CREATE TABLE records (year STRING, temperature INT, quality INT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t';
```

# Data Model - Tables

- Tables
  - Analogous to **tables** in relational DBs.
  - Each table has a corresponding **directory** in HDFS.
  - Creating a table

CREATE TABLE records (year STRING, temperature INT, quality INT)

**ROW FORMAT DELIMITED**

**FIELDS TERMINATED BY '\t';**

- Loading a table

LOAD DATA INPATH 'input/ncdc/micro-tab/sample.txt'

OVERWRITE INTO TABLE records;

# Data Model - Tables

- Location: tables are stored as **directories** under Hive's warehouse
  - Default **/user/hive/warehouse**
  - Controlled by the *hive.metastore.warehouse.dir* property

```
$ ... ls /user/hive/warehouse/records/  
sample.txt
```

- Query: once loaded, we can run a query against it:  
hive> **SELECT year, MAX(temperature)**  
    > **FROM records**  
    > **WHERE temperature != 9999 AND quality IN (0, 1, 4, 5, 9)**  
    > **GROUP BY year;**

```
1949 111
```

```
1950 22
```

# Managed vs. External Tables

- Managed Tables
  - **Default** option when a table is created
  - Hive **moves** the data into its **warehouse** directory

```
CREATE TABLE managed_table (dummy STRING);
```

```
LOAD DATA INPATH '/user/tom/data.txt' INTO table managed_table;
```

- Drop table semantics: the table, including its **metadata** *and its data*, is deleted

```
DROP TABLE managed_table;
```

# Managed vs. External Tables

- External Tables
  - Point to **existing** data directories in HDFS
  - Data is assumed to be in Hive-compatible format

```
CREATE EXTERNAL TABLE external_table (dummy STRING)  
LOCATION '/user/tom/external_table';
```

```
LOAD DATA INPATH '/user/tom/data.txt' INTO TABLE external_table;
```

- Drop table semantics: drops **only** the **metadata**, leave the data untouched

# Table - Partitions

- Partitions
  - A way of dividing a table into **coarse-grained** parts
  - Can make it **faster** to do queries on slices of the data
  - A table may be partitioned in multiple dimensions
  - Example
    - Table: a log file; Partition columns: (datestamp, country)

```
/user/hive/warehouse/logs
```

```
|— dt=2001-01-01/  
|   |— country=GB/  
|   |   |— file1  
|   |   |— file2  
|   |— country=US/  
|   |   |— file3  
|— dt=2001-01-02/  
|   |— country=GB/  
|   |   |— file4  
|   |— country=US/  
|   |   |— file5  
|   |   |— file6
```

# Table - Partitions

- Creating Partitions

- Defined at table creation time using the PARTITIONED BY clause

```
CREATE TABLE logs (ts BIGINT, line STRING)  
PARTITIONED BY (dt STRING, country STRING);
```

- While loading data, the partition values are specified explicitly:

```
LOAD DATA LOCAL INPATH 'input/hive/partitions/file1'  
INTO TABLE logs  
PARTITION (dt='2001-01-01', country='GB');
```

# Table - Partitions

- Ask Hive to show Partitions  
`hive> SHOW PARTITIONS logs;`

dt=2001-01-01/country=GB

dt=2001-01-01/country=US

dt=2001-01-02/country=GB

dt=2001-01-02/country=US


- Query on specific Partitions  
`SELECT ts, dt, line FROM logs  
WHERE country='GB';`

will only scan *file1*, *file2*, and *file4*

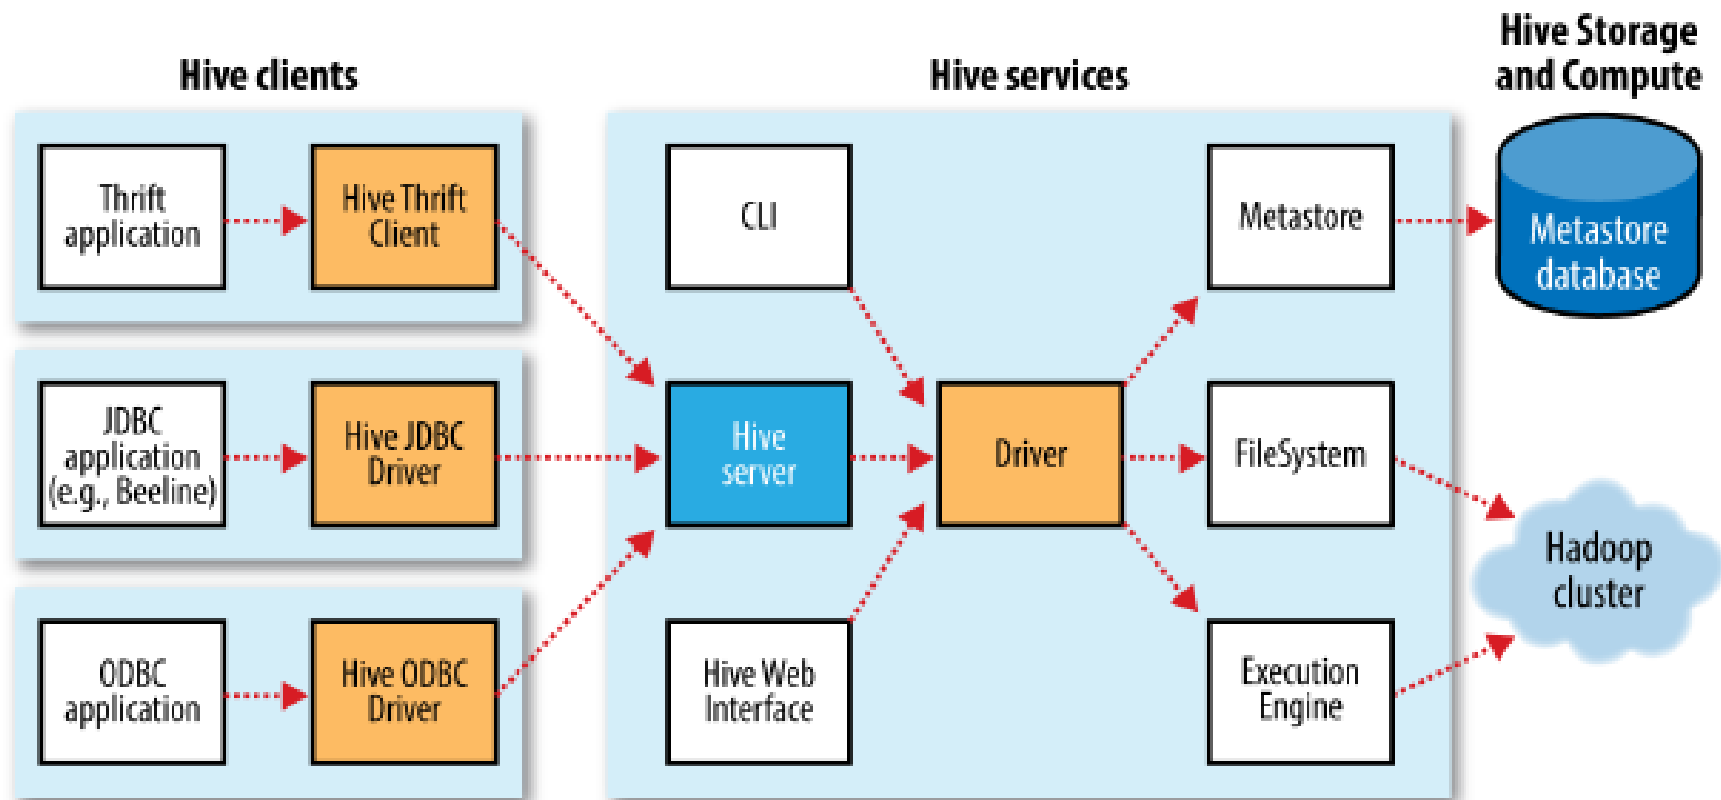




# Outline

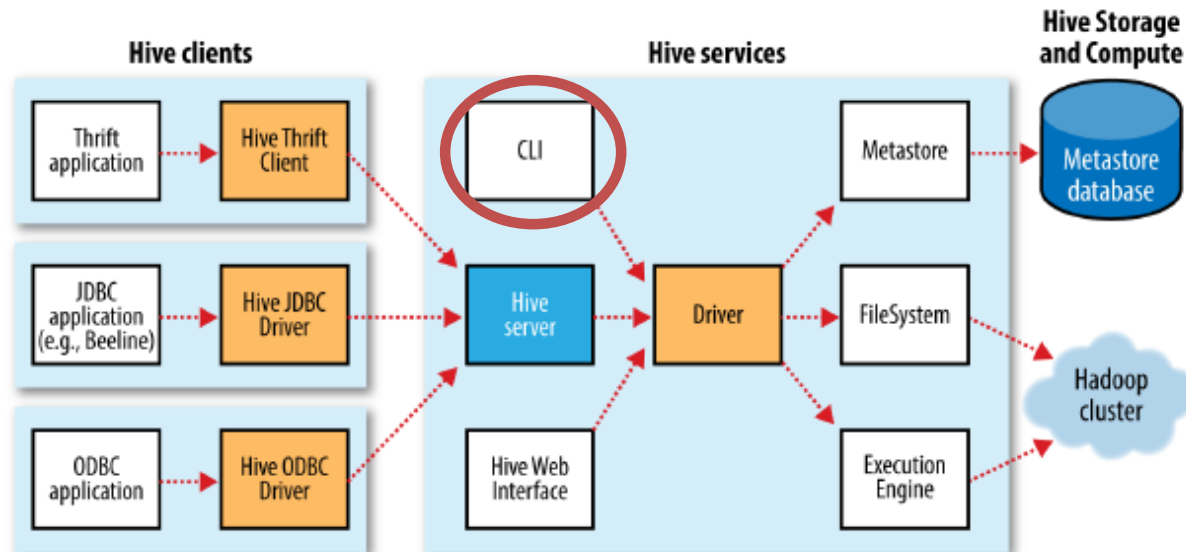
- Hive
  - Data model
  - System architecture 
  - HiveQL
  - File formats
- File-based Data Structures
  - Row-oriented storage formats
  - Column-oriented storage formats

# Hive Architecture



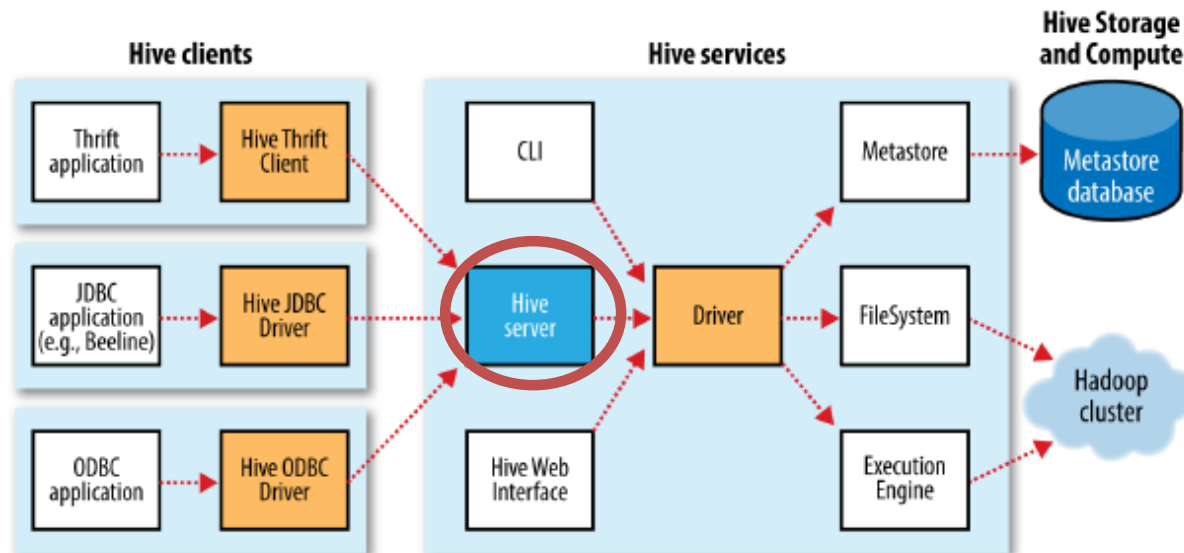
# Hive Services

- cli
  - The **command-line** interface to Hive (the shell). This is the **default** service.




# Hive Services

- hiveserver2
  - Runs Hive as a **server**, enabling access from **a range of clients**
  - Applications using the Thrift, JDBC, and ODBC connectors need to run a Hive server to communicate with Hive



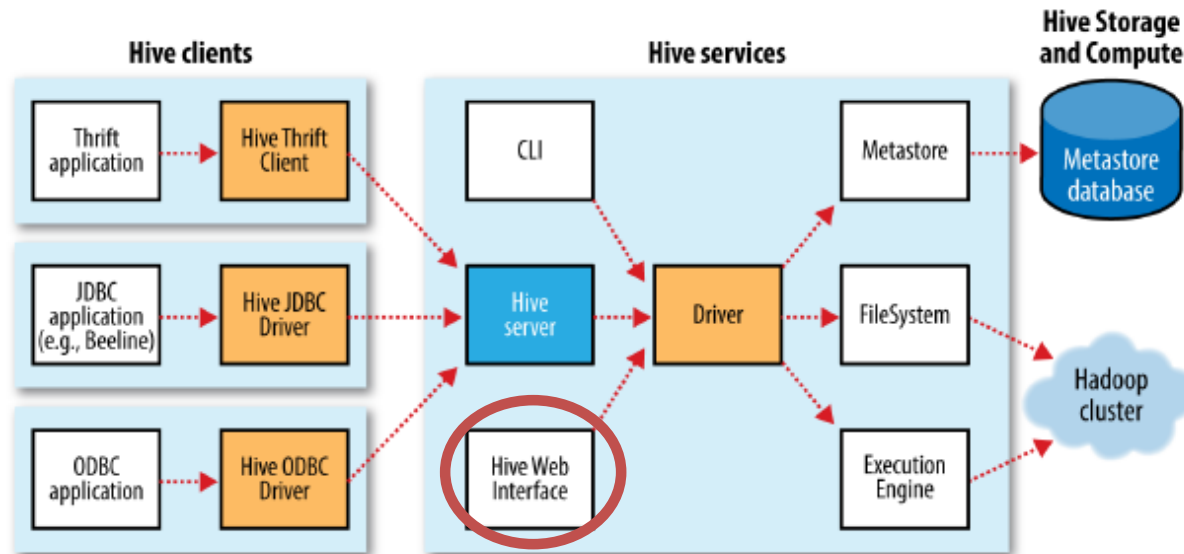
# Hive Services

- metastore
  - The **central repository** of Hive metadata



# Hive Services

- hwi
  - A simple **web interface** that can be used as an alternative to the CLI without having to install any client software




# Hive Services

- beeline
  - A **command-line** interface to Hive that works in **embedded** mode, or by **connecting to a hiveserver2** process using **JDBC**

# Metastore configurations


- Embedded metastore
  - The metastore service runs in the **same JVM** as the Hive service and contains an embedded **Derby database** instance backed by the local disk.
  - Can have only **one** Hive session open at a time that accesses the same metastore






# Metastore configurations

- Local metastore
  - Supports **multiple sessions** (and therefore multiple users)
  - Metastore service still runs in the **same process** as the Hive service, but connects to a **database** running in a **separate process**
  - **MySQL** is a popular choice




# Metastore configurations

- Remote metastore
  - One or more metastore servers run in **separate processes** to the Hive service
  - Brings better manageability and security because the database tier can be completely firewalled off



# Outline

- Hive
  - Data model
  - System architecture
  - HiveQL 
  - File formats
- File-based Data Structures
  - Row-oriented storage formats
  - Column-oriented storage formats

# Hive vs. traditional relational databases

	Hive	Relational databases
Schema on Read Versus Schema on Write	<b>Schema on read.</b> Does not verify the schema when the data is loaded, but rather when the a query is issued	<b>Schema on write.</b> A table's schema is <b>enforced</b> at data load time. If the data being loaded doesn't conform to the schema, then it is rejected.
In-place updates	<b>Not supported.</b> Changes resulting from inserts, updates, and deletes are <b>stored</b> in <b>small delta files</b> . They are <b>periodically merged</b> into the base table files by MapReduce jobs that are run in the background	<b>Supported.</b> Updates are reflected in <b>real-time</b> (mostly).
Index	There are currently <b>two</b> index types: <i>compact</i> and <i>bitmap</i> . (no longer supported since v3)	Various sophisticated indexes supported

# Hive Query Language (HiveQL)

- Basic SQL
  - From clause sub-query
  - ANSI JOIN (theta join supported since 0.14)
  - Multi group-by
  - Sampling
  - Objects Traversal
- Extensibility
  - Pluggable Map-reduce scripts using TRANSFORM

# Type System

- Primitive types
  - Integers: TINYINT, SMALLINT, INT, BIGINT.
  - Boolean: BOOLEAN.
  - Floating point numbers: FLOAT, DOUBLE .
  - String: STRING.
- Complex types
  - Structs: {a INT; b INT}.
  - Maps: M['group'].
  - Arrays: ['a', 'b', 'c'], A[1] returns 'b'.

# Complex type example

- Create table

```
create table tab11
```

```
(id int,
```

```
name string,
```

```
sal bigint,
```

```
sub array<string>,
```

```
dud map<string,int>,
```

```
addr struct<city:string,state:string,pin:bigint>
```

```
)
```

```
row format delimited
```

```
fields terminated by ','
```

```
collection items terminated by '$'
```

```
map keys terminated by '#';
```

- Data file

```
1,abc,40000,a$b$c,pf#500$epf#200,hyd$ap$500001
```

```
2,def,3000,d$f,pf#500,bang$kar$600038
```

# Hive Query Language (HiveQL) - contd.

- Insertion

```
INSERT OVERWRITE TABLE sample1 '/tmp/hdfs_out' SELECT * FROM sample  
WHERE ds='2012-02-24';
```

```
INSERT OVERWRITE DIRECTORY '/tmp/hdfs_out' SELECT * FROM sample  
WHERE ds='2012-02-24';
```

```
INSERT OVERWRITE LOCAL DIRECTORY '/tmp/hive-sample-out' SELECT *  
FROM sample;
```



# HiveQL – Join

- Hive equi-join

page\_view

pageid	<b>userid</b>	time
1	<b>111</b>	9:08:01
2	<b>111</b>	9:08:13
1	<b>222</b>	9:08:14

X

user

<b>userid</b>	age	gender
<b>111</b>	25	female
<b>222</b>	32	male

=

pv\_users

pageid	age
1	25
2	25
1	32

- SQL:

```
INSERT INTO TABLE pv_users
```

```
SELECT pv.pageid, u.age
```

```
FROM page_view pv JOIN user u ON (pv.userid =  
u.userid);
```

# HiveQL – Join in Map Reduce

page\_view

pageid	userid	time
1	111	9:08:01
2	111	9:08:13
1	222	9:08:14

Map

key	value
111	<1,1>
111	<1,2>
222	<1,1>

user

userid	age	gender
111	25	female
222	32	male

key	value
111	<2,25>
222	<2,32>

Shuffle  
Sort

key	value
111	<1,1>
111	<1,2>
111	<2,25>

key	value
222	<1,1>
222	<2,32>


pv\_users

pageid	age
1	25
2	25

Reduce

pageid	age
1	32

# HiveQL – Group By




- SQL:

```
INSERT INTO TABLE pageid_age_sum
SELECT pageid, age, count(1)
FROM pv_users
GROUP BY pageid, age;
```

# HiveQL – Group By in Map Reduce


pv\_users

pageid	age
1	25
2	25



key	value
<1,25>	1
<2,25>	1

Map




pageid	age
1	32
2	25

key	value
<1,32>	1
<2,25>	1

Shuffle  
Sort

key	value
<1,25>	1
<1,32>	1



Reduce



key	value
<2,25>	1
<2,25>	1

pageid\_age\_sum

pageid	age	Count
1	25	1
1	32	1

pageid	age	Count
2	25	2